

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("C:/Users/abds0/Desktop/DAP project/dataset.csv")
print("Rows and Columns:", df.shape)
print("\nFirst 5 rows:")
print(df.head())
categorical_cols = [
    "Attrition", "BusinessTravel", "Department", "Gender",
    "JobRole", "MaritalStatus"]
for col in categorical_cols:
    print(f"\nValue counts for {col}:")
    print(df[col].value_counts())

```

Rows and Columns: (1470, 31)

First 5 rows:

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102	Sales	
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	

	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	\
0	1	2	Life Sciences	2	
1	8	1	Life Sciences	3	
2	2	2	Other	4	
3	3	4	Life Sciences	4	
4	2	1	Medical	1	

	Gender	...	PerformanceRating	RelationshipSatisfaction	StockOptionLevel	\
0	Female	...	3	1	0	
1	Male	...	4	4	1	
2	Male	...	3	2	0	
3	Female	...	3	3	0	
4	Male	...	3	4	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager	
0	4	0	5	
1	7	1	7	
2	0	0	0	
3	7	3	0	
4	2	2	2	

[5 rows x 31 columns]

```

Value counts for Attrition:
Attrition
No      1233
Yes     237
Name: count, dtype: int64

Value counts for Department:
Department
Research & Development    961
Sales                  446
Human Resources        63
Name: count, dtype: int64

Value counts for Gender:
Gender
Male      882
Female    588
Name: count, dtype: int64

Value counts for JobRole:
JobRole
Sales Executive          326
Research Scientist        292
Laboratory Technician     259
Manufacturing Director    145
Healthcare Representative 131
Manager                 102
Sales Representative      83
Research Director         80
Human Resources           52
Name: count, dtype: int64

Value counts for MaritalStatus:
MaritalStatus
Married      673
Single       470
Divorced     327
Name: count, dtype: int64

import pandas as pd
df = pd.read_csv("C:/Users/abds0/Desktop/DAP project/dataset.csv")
print("\nMissing values in each column:")
print(df.isnull().sum())
print("\nDuplicate rows before:", df.duplicated().sum())
df = df.drop_duplicates()
print("Duplicate rows after:", df.duplicated().sum())
df['OverTime'] = df['OverTime'].str.strip()
df['OverTime'] = df['OverTime'].replace({
    'yes': 'Yes',
    'YES': 'Yes',
    'no': 'No',
    'NO': 'No'
})
remove_cols = ["EmployeeCount", "Over18", "StandardHours"]
df = df.drop(columns=["EmployeeCount", "Over18", "StandardHours"], errors='ignore')
print("\nDropped columns:", remove_cols)
df['Attrition'] = df['Attrition'].astype('category')
df['Gender'] = df['Gender'].astype('category')
print("\nUpdated column types:")
print(df.dtypes)

```

```

Missing values in each column:
Age                      0
Attrition                 0
BusinessTravel              0
DailyRate                  0
Department                 0
DistanceFromHome             0
Education                  0
EducationField                0
EnvironmentSatisfaction        0
Gender                     0
HourlyRate                  0
JobInvolvement                0
JobLevel                    0
JobRole                     0
JobSatisfaction                0
MaritalStatus                 0
MonthlyIncome                 0
MonthlyRate                  0
NumCompaniesWorked             0
OverTime                     0
PercentSalaryHike               0
PerformanceRating               0
RelationshipSatisfaction        0
StockOptionLevel                0
TotalWorkingYears                0
TrainingTimesLastYear             0
WorkLifeBalance                 0
YearsAtCompany                  0
YearsInCurrentRole                0
YearsSinceLastPromotion             0
YearsWithCurrManager                0
dtype: int64

Dropped columns: ['EmployeeCount', 'Over18', 'StandardHours']

Updated column types:
Age                         int64
Attrition                   category
BusinessTravel                object
DailyRate                     int64
Department                   object
DistanceFromHome                int64
Education                     int64
EducationField                  object
EnvironmentSatisfaction        int64
Gender                       category
HourlyRate                     int64
JobInvolvement                  int64
JobLevel                      int64
JobRole                        object
JobSatisfaction                  int64
MaritalStatus                   object
MonthlyIncome                   int64
MonthlyRate                     int64
NumCompaniesWorked                int64
OverTime                      object
PercentSalaryHike                  int64
PerformanceRating                  int64
RelationshipSatisfaction        int64
StockOptionLevel                  int64
TotalWorkingYears                  int64
TrainingTimesLastYear                int64
WorkLifeBalance                  int64
YearsAtCompany                   int64
YearsInCurrentRole                  int64
YearsSinceLastPromotion                int64
YearsWithCurrManager                  int64
dtype: object


```


Duplicate rows before: 0
 Duplicate rows after: 0

```

import matplotlib.pyplot as plt
import seaborn as sns
import math
num_cols = [
    "Age", "DailyRate", "DistanceFromHome", "MonthlyIncome",
    "MonthlyRate", "NumCompaniesWorked", "TotalWorkingYears",
    "TrainingTimesLastYear", "YearsAtCompany", "YearsInCurrentRole",
    "YearsSinceLastPromotion", "YearsWithCurrManager"]
num_cols = [c for c in num_cols if c in df.columns]
n = len(num_cols)
rows = math.ceil(n / 3) # grid with 3 columns
cols = 3
plt.figure(figsize=(15, 4 * rows))
for i, col in enumerate(num_cols, 1):
    plt.subplot(rows, cols, i)
    sns.boxplot(x=df[col], width=0.4)
    plt.title(col, fontsize=10)
    plt.tight_layout()
plt.suptitle("Boxplots of Numerical Features (Compact View)", fontsize=14, y=1.02)
plt.show()

```





```

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(14, 12))

# ----- Attrition vs Age -----
plt.subplot(3, 2, 1)
sns.boxplot(x='Attrition', y='Age', data=df)
plt.title("Attrition vs Age")

# ----- Attrition vs MonthlyIncome -----
plt.subplot(3, 2, 2)
sns.boxplot(x='Attrition', y='MonthlyIncome', data=df)
plt.title("Attrition vs Monthly Income")

# ----- Attrition vs OverTime -----
plt.subplot(3, 2, 3)
sns.countplot(x='OverTime', hue='Attrition', data=df)
plt.title("Attrition by OverTime")
plt.xticks(rotation=30)

# ----- Attrition vs JobSatisfaction -----
plt.subplot(3, 2, 4)
sns.boxplot(x='Attrition', y='JobSatisfaction', data=df)
plt.title("Attrition vs Job Satisfaction")

plt.tight_layout()
plt.show()

```





```

# Select important numerical columns for Pearson
pearson_cols = [
    'Age', 'MonthlyIncome', 'DistanceFromHome',
    'DailyRate', 'MonthlyRate',
    'TotalWorkingYears', 'YearsAtCompany'
]

print("\n✿ Pearson Correlation Matrix:")
pearson_corr = df[pearson_cols].corr(method='pearson')
print(pearson_corr)


# Heatmap
plt.figure(figsize=(10,6))
sns.heatmap(pearson_corr, annot=True, cmap='coolwarm')
plt.title("Pearson Correlation Heatmap")
plt.show()

```



❖ Pearson Correlation Matrix:

	Age	MonthlyIncome	DistanceFromHome	DailyRate	\	
Age	1.000000	0.497855	-0.001686	0.010661		
MonthlyIncome	0.497855	1.000000	-0.017014	0.007707		
DistanceFromHome	-0.001686	-0.017014	1.000000	-0.004985		
DailyRate	0.010661	0.007707	-0.004985	1.000000		
MonthlyRate	0.028051	0.034814	0.027473	-0.032182		
TotalWorkingYears	0.680381	0.772893	0.004628	0.014515		
YearsAtCompany	0.311309	0.514285	0.009508	-0.034055		
		MonthlyRate	TotalWorkingYears	YearsAtCompany		
Age		0.028051	0.680381	0.311309		
MonthlyIncome		0.034814	0.772893	0.514285		
DistanceFromHome		0.027473	0.004628	0.009508		
DailyRate		-0.032182	0.014515	-0.034055		
MonthlyRate		1.000000	0.026442	-0.023655		
TotalWorkingYears		0.026442	1.000000	0.628133		
YearsAtCompany		-0.023655	0.628133	1.000000		



```
spearman_cols = [
    'Education', 'JobLevel', 'WorkLifeBalance',
    'JobSatisfaction', 'EnvironmentSatisfaction'
]

print("\n❖ Spearman Correlation Matrix:")
spearman_corr = df[spearman_cols].corr(method='spearman')
print(spearman_corr)

# Heatmap
plt.figure(figsize=(8,5))
sns.heatmap(spearman_corr, annot=True, cmap='viridis')
plt.title("Spearman Correlation Heatmap (Ordinal Columns)")
plt.show()
```

❖ Spearman Correlation Matrix:


```

Education      1.000000  0.107419   0.017350 \\
JobLevel       0.107419  1.000000   0.040466 \\
WorkLifeBalance 0.017350  0.040466  1.000000 \\
JobSatisfaction -0.005175 -0.000852  -0.029781 \\
EnvironmentSatisfaction -0.027625 -0.000192   0.027169 \\

                                         JobSatisfaction EnvironmentSatisfaction
Education      -0.005175      -0.027625 \\
JobLevel       -0.000852      -0.000192 \\
WorkLifeBalance -0.029781      0.027169 \\
JobSatisfaction 1.000000     -0.002993 \\
EnvironmentSatisfaction -0.002993      1.000000 \\

```

Spearman Correlation Heatmap (Ordinal Columns)



```

df['Attrition'] = df['Attrition'].astype(str).str.strip().str.title()
df['Attrition_numeric'] = df['Attrition'].map({'Yes': 1, 'No': 0})
print("\n❖ Correlation of Numerical Features with Attrition:")
corr_attr = df.corr(numeric_only=True)[['Attrition_numeric']].sort_values(ascending=False)
print(corr_attr)

```

❖ Correlation of Numerical Features with Attrition:

Attrition_numeric	1.000000
DistanceFromHome	0.077924
NumCompaniesWorked	0.043494
MonthlyRate	0.015170
PerformanceRating	0.002889
HourlyRate	-0.006846
PercentSalaryHike	-0.013478
Education	-0.031373
YearsSinceLastPromotion	-0.033019
RelationshipSatisfaction	-0.045872
DailyRate	-0.056652
TrainingTimesLastYear	-0.059478
WorkLifeBalance	-0.063939
EnvironmentSatisfaction	-0.103369
JobSatisfaction	-0.103481
JobInvolvement	-0.130016
YearsAtCompany	-0.134392
StockOptionLevel	-0.137145
YearsWithCurrManager	-0.156199
Age	-0.159205
MonthlyIncome	-0.159840
YearsInCurrentRole	-0.160545
JobLevel	-0.169105
TotalWorkingYears	-0.171063

Name: Attrition_numeric, dtype: float64

- ◆ CHI-SQUARE TESTS (Categorical vs Categorical)

```
Gender: p-value = 0.2906 --> Not Significant
Department: p-value = 0.0045 --> Significant
BusinessTravel: p-value = 0.0000 --> Significant
JobRole: p-value = 0.0000 --> Significant
MaritalStatus: p-value = 0.0000 --> Significant
EducationField: p-value = 0.0068 --> Significant
OverTime: p-value = 0.0000 --> Significant
```

```
from scipy.stats import ttest_ind

print("\n\n♦ t-TESTS (Numerical vs Attrition)\n")

df['Attrition_numeric'] = df['Attrition'].map({'Yes':1, 'No':0})

num_cols_to_test = [
    'MonthlyIncome', 'Age', 'DailyRate', 'HourlyRate',
    'MonthlyRate', 'TotalWorkingYears', 'DistanceFromHome',
    'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion'
]

for col in num_cols_to_test:
    yes = df[df['Attrition']=='Yes'][col].dropna()
    no = df[df['Attrition']=='No'][col].dropna()

    t_stat, p = ttest_ind(yes, no, equal_var=False)

    print(f'{col}: p-value = {p:.4f} --> {"Significant" if p < 0.05 else "Not Significant"}')
```

- ◆ t-TESTS (Numerical vs Attrition)

```
MonthlyIncome: p-value = 0.0000 --> Significant
Age: p-value = 0.0000 --> Significant
DailyRate: p-value = 0.0300 --> Significant
HourlyRate: p-value = 0.7914 --> Not Significant
MonthlyRate: p-value = 0.5653 --> Not Significant
TotalWorkingYears: p-value = 0.0000 --> Significant
DistanceFromHome: p-value = 0.0041 --> Significant
YearsAtCompany: p-value = 0.0000 --> Significant
YearsInCurrentRole: p-value = 0.0000 --> Significant
YearsSinceLastPromotion: p-value = 0.1987 --> Not Significant
```

```
from scipy.stats import chi2_contingency
import pandas as pd

print("\n♦ CHI-SQUARE TESTS (Categorical vs Categorical)\n")

cat_cols = [
    'Gender', 'Department', 'BusinessTravel', 'JobRole',
    'MaritalStatus', 'EducationField', 'OverTime'
]

chi_results = {}

for col in cat_cols:
    table = pd.crosstab(df['Attrition'], df[col])
    chi2, p, dof, exp = chi2_contingency(table)
    chi_results[col] = p
    print(f'{col}: p-value = {p:.4f} --> {"Significant" if p < 0.05 else "Not Significant"}')
```

CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.87	1.00	0.93	255
1	0.50	0.03	0.05	39
accuracy			0.87	294
macro avg	0.68	0.51	0.49	294
weighted avg	0.82	0.87	0.81	294

Model is trained successfully!

```
sample = pd.DataFrame({
    "Age": [22],
    "MonthlyIncome": [5000],
    "OverTime": ["Yes"],
    "JobSatisfaction": [1],
    "DistanceFromHome": [25]
})
prediction = model.predict(sample)[0]
prob = model.predict_proba(sample)[0][1]
print("Prediction =", prediction)
print("Probability of Attrition =", prob)
```

Prediction = 1
Probability of Attrition = 0.6709265976809173

```
new_data = pd.DataFrame({
    "Age": [25, 18],
    "MonthlyIncome": [20000, 7000],
    "OverTime": ["No", "Yes"],
    "JobSatisfaction": [4, 2],
    "DistanceFromHome": [5, 13]
})
preds = model.predict(new_data)
probs = model.predict_proba(new_data)[:,1]
print(preds)
print(probs)
```

[0 1]
[0.01970809 0.50760877]

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

df = pd.read_csv("C:/Users/abds0/Desktop/DAP project/dataset.csv")
df["Attrition"] = df["Attrition"].map({"Yes": 1, "No": 0})
simple_features = ["Age", "MonthlyIncome", "OverTime", "JobSatisfaction", "DistanceFromHome"]
X = df[simple_features]
y = df["Attrition"]
cat_cols = ["OverTime"]
num_cols = ["Age", "MonthlyIncome", "JobSatisfaction", "DistanceFromHome"]
preprocess = ColumnTransformer([
    ("cat", OneHotEncoder(drop='first'), cat_cols),
    ("num", "passthrough", num_cols)])
model = Pipeline([
    ("preprocess", preprocess),
    ("classifier", LogisticRegression(max_iter=300))])
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("\nCLASSIFICATION REPORT:")
print(classification_report(y_test, y_pred))
print("\nModel is trained successfully!")
```